

```

1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 import math
5 import numpy as np
6 import torch.cuda.amp as amp
7 from tqdm import tqdm
8 import matplotlib.pyplot as plt
9 from torch.utils.data import DataLoader, TensorDataset

1 # Training Configuration
2 epochs = 5
3 batch_size = 1024
4 lr = 3e-4
5 weight_decay = 0.01
6 device = "cuda"
7 checkpoint_filepath = None # Set to a path if you want to load a checkpoint
8 save_dir = "checkpoints"
9 import os
10 os.makedirs(save_dir, exist_ok=True)

```

✓ Model code

```

1 class PatchEmbedding(nn.Module):
2     """
3     Module that converts image patches to embeddings for Vision Transformer.
4     """
5     def __init__(self,
6                 image_size: tuple = (64, 72),
7                 patch_size: int = 8,
8                 in_channels: int = 3,
9                 embedding_dim: int = 1024):
10         super().__init__()
11         self.image_size = image_size
12         self.patch_size = patch_size
13         self.in_channels = in_channels
14
15         # Calculate number of patches
16         self.num_patches = (image_size[0] // patch_size) * (image_size[1] // patch_size)
17

```

```

18     # Create projection for converting patches to embeddings
19     self.projection = nn.Conv2d(
20         in_channels=in_channels,
21         out_channels=embedding_dim,
22         kernel_size=patch_size,
23         stride=patch_size
24     )
25
26     # CLS token embedding
27     self.cls_token = nn.Parameter(torch.zeros(1, 1, embedding_dim))
28
29     # Positional embedding (Normal distribution initialization of value)
30     self.positions = nn.Parameter(torch.zeros(1, self.num_patches + 1, embedding_dim))
31     nn.init.trunc_normal_(self.positions, std=0.02)
32
33     def forward(self, x: torch.Tensor) -> torch.Tensor:
34         batch_size = x.shape[0]
35
36         # Convert image to patches and project to embedding dimension
37         # x shape: [batch_size, channels, height, width]
38         x = self.projection(x)
39         # x shape: [batch_size, embedding_dim, height/patch_size, width/patch_size]
40
41         # Flatten patches to sequence
42         x = x.flatten(2).transpose(1, 2)
43         # x shape: [batch_size, num_patches, embedding_dim]
44
45         # Add CLS token
46         cls_tokens = self.cls_token.expand(batch_size, -1, -1)
47         x = torch.cat((cls_tokens, x), dim=1)
48
49         # Add positional embeddings
50         x = x + self.positions
51
52         return x
53
54
55 class VisionAttention(nn.Module):
56     def __init__(self,
57         hidden_dim: int,
58         head_dim: int,
59         q_head: int,
60         kv_head: int,
61         lora_rank: int = 16):
62         super().__init__()

```

```

63     self.head_dim = head_dim
64     self.q_head = q_head
65     self.kv_head = kv_head
66     self.qkv = nn.Linear(hidden_dim, (q_head+kv_head*2)*head_dim)
67     self.o = nn.Linear(q_head*head_dim, hidden_dim)
68     self.scaler = 1/math.sqrt(head_dim)
69     self.lora_qkv_a = nn.Linear(hidden_dim, lora_rank)
70     self.lora_qkv_b = nn.Linear(lora_rank, (q_head+kv_head*2)*head_dim)
71     self.lora_o_a = nn.Linear(q_head*head_dim, lora_rank)
72     self.lora_o_b = nn.Linear(lora_rank, hidden_dim)
73
74     if q_head != kv_head:
75         # If we are using multi query attention
76         assert q_head % kv_head == 0
77         self.multi_query_attention = True
78         self.q_kv_scale = q_head//kv_head
79     else:
80         self.multi_query_attention = False
81
82     def forward(self, tensor: torch.Tensor, attention_mask: torch.Tensor = None):
83         batch_size, seq_len, hid_dim = tensor.shape
84
85         qkv_tensor = self.qkv(tensor)
86         if fine_tuning:
87             lora_tensor = self.lora_qkv_a(tensor)
88             lora_tensor = self.lora_qkv_b(lora_tensor)
89             qkv_tensor = lora_tensor + qkv_tensor
90         query, key, value = qkv_tensor.split([self.head_dim*self.q_head, self.head_dim*self.kv_head*2], dim=-1)
91
92         query = query.view(batch_size, seq_len, self.q_head, self.head_dim)
93         key = key.view(batch_size, seq_len, self.kv_head, self.head_dim)
94         value = value.view(batch_size, seq_len, self.kv_head, self.head_dim)
95
96         if self.multi_query_attention:
97             # If we are using multi query attention, duplicate key value head:
98             key = torch.repeat_interleave(key, self.q_kv_scale, dim=-2)
99             value = torch.repeat_interleave(value, self.q_kv_scale, dim=-2)
100
101         # Switch to batch_size, head, seq_len, head_dim
102         query = query.transpose(1, 2)
103         key = key.transpose(1, 2)
104         value = value.transpose(1, 2)
105
106         # Classic self attention
107         attention_raw = torch.matmul(query, key.transpose(2, 3))
108         attention = torch.softmax(attention_raw, dim=-1)
109         attention = attention * attention_mask.unsqueeze(-1)
110         output = torch.matmul(attention, value)
111         output = self.o(output)
112         return output

```

```

108         attention_scaled = attention_raw * self.scaler
109         if attention_mask != None:
110             attention_scaled += attention_mask
111         attention_score = torch.softmax(attention_scaled, dim=-1)
112         value = torch.matmul(attention_score, value)
113
114         # Reshape back to batch_size, seq_len, hid_dim
115         value = value.transpose(1, 2).contiguous()
116         value = value.view(batch_size, seq_len, hid_dim)
117
118         # Output layer
119         output = self.o(value)
120         if fine_tuning:
121             lora_tensor = self.lora_o_a(value)
122             lora_tensor = self.lora_o_b(lora_tensor)
123             output = lora_tensor + output
124
125         return output
126
127
128 class FeedForward(nn.Module):
129     def __init__(self,
130                 hidden_size: int,
131                 expansion_factor: int = 4,
132                 dropout_ratio: float = 0.1,
133                 lora_rank: int = 16):
134         super().__init__()
135         self.gate_and_up = nn.Linear(hidden_size, hidden_size * expansion_factor)
136         self.down = nn.Linear(hidden_size * expansion_factor, hidden_size)
137         self.dropout = nn.Dropout(p=dropout_ratio)
138         self.lora_gate_and_up_a = nn.Linear(hidden_size, lora_rank)
139         self.lora_gate_and_up_b = nn.Linear(lora_rank, hidden_size * expansion_factor)
140         self.lora_down_a = nn.Linear(hidden_size * expansion_factor, lora_rank)
141         self.lora_down_b = nn.Linear(lora_rank, hidden_size)
142
143     def forward(self, tensor: torch.Tensor, fine_tuning: bool = False) -> torch.Tensor:
144         gate_and_up = self.gate_and_up(tensor)
145         if fine_tuning:
146             lora_tensor = self.lora_gate_and_up_a(tensor)
147             lora_tensor = self.lora_gate_and_up_b(lora_tensor)
148             gate_and_up = gate_and_up + lora_tensor
149         gate, up = gate_and_up.chunk(chunks=2, dim=-1)
150         gate = F.gelu(gate, approximate="tanh")
151         tensor = gate * up
152         tensor = self.dropout(tensor)
153         down_tensor = self.down(tensor)

```

```

153         down_tensor = self.down(tensor)
154     if fine_tuning:
155         lora_tensor = self.lora_down_a(tensor)
156         lora_tensor = self.lora_down_b(lora_tensor)
157         down_tensor = down_tensor + lora_tensor
158     return down_tensor
159
160
161 class MOE(nn.Module):
162     def __init__(self, hidden_size: int, device: str, num_experts: int = 8, expansion_factor: int = 2):
163         super().__init__()
164         self.gate = nn.Linear(hidden_size, num_experts)
165         self.num_experts = num_experts
166         self.device = device
167         self.experts = nn.ModuleList([FeedForward(hidden_size, expansion_factor) for _ in range(num_experts)])
168
169     def forward(self, tensor: torch.Tensor, fine_tuning: bool = False) -> tuple:
170         # Flatten for better manipulation, this is ok because tokens are independent
171         batch_size, seq_len, hidden_size = tensor.shape
172         flat_tensor = tensor.reshape(batch_size * seq_len, hidden_size)
173
174         # Pass through the gating network and select experts
175         tensor = self.gate(flat_tensor)
176         tensor = F.softmax(tensor, dim=-1)
177
178         # The output of this step is a tensor of shape [batch_size * seq_len, num_experts]
179         value_tensor, index_tensor = tensor.topk(k=2, dim=-1)
180
181         # Find the load balancing loss
182         counts = torch.bincount(index_tensor[:, 0], minlength=self.num_experts)
183         frequencies = counts.float() / (batch_size * seq_len) # This is the hard probability
184         probability = tensor.mean(0) # This is the soft probability
185         load_balancing_loss = (probability * frequencies).mean() * float(self.num_experts)
186
187         # Normalize top1 and top2 score
188         top_expert_score = value_tensor[:, 0]
189         second_expert_score = value_tensor[:, 1]
190         total_score = top_expert_score + second_expert_score
191         top_expert_score = top_expert_score / total_score
192         second_expert_score = second_expert_score / total_score
193
194         # Split into top 2 experts
195         split_tensors = torch.split(index_tensor, 1, dim=-1)
196         top_expert, second_expert = split_tensors[0], split_tensors[1]
197         indices = torch.arange(batch_size * seq_len).unsqueeze(-1).to(self.device)
198         top_expert = torch.cat((indices * top_expert, dim=-1)

```

```

198 top_expert = torch.cat((indices, top_expert), dim=-1)
199 second_expert = torch.cat((indices, second_expert), dim=-1)
200
201 # Sort based on expert selection
202 top_expert = top_expert[top_expert[:,1].argsort()]
203 second_expert = second_expert[second_expert[:,1].argsort()]
204
205 # Count how many tokens goes to each expert
206 top_expert_counts = torch.zeros(self.num_experts, dtype=int)
207 for i in range(self.num_experts):
208     top_expert_counts[i] = (top_expert[:,1] == i).sum()
209 top_expert_counts = top_expert_counts.tolist()
210
211 second_expert_counts = torch.zeros(self.num_experts, dtype=int)
212 for i in range(self.num_experts):
213     second_expert_counts[i] = (second_expert[:,1] == i).sum()
214 second_expert_counts = second_expert_counts.tolist()
215
216 # Split input tokens for each expert
217 top_expert_tokens = flat_tensor[top_expert[:,0]]
218 second_expert_tokens = flat_tensor[second_expert[:,0]]
219
220 # Split into a list of tensors, element i tensor is for ith expert.
221 top_expert_tokens = torch.split(top_expert_tokens, top_expert_counts,
222 second_expert_tokens = torch.split(second_expert_tokens, second_expert_counts,
223
224 # Input into each expert and obtain results in a list
225 top_expert_outputs = [self.experts[i](top_expert_tokens[i], fine_tune)
226 second_expert_outputs = [self.experts[i](second_expert_tokens[i], fine_tune)
227
228 # Combine outputs
229 top_expert_outputs = torch.cat(top_expert_outputs, dim=0)
230 second_expert_outputs = torch.cat(second_expert_outputs, dim=0)
231
232 # Re-index the output back to original token order
233 flat_top_expert_tensor = torch.zeros_like(flat_tensor, dtype=torch.float)
234 flat_top_expert_tensor.index_copy_(0, top_expert[:, 0], top_expert_outputs)
235
236 flat_second_expert_tensor = torch.zeros_like(flat_tensor, dtype=torch.float)
237 flat_second_expert_tensor.index_copy_(0, second_expert[:, 0], second_expert_outputs)
238
239 # Find final output tensor based on weight between top and second expert
240 final_tensor = top_expert_score.unsqueeze(-1) * flat_top_expert_tensor +
241 second_expert_score.unsqueeze(-1) * flat_second_expert_tensor
242
243 # Reshape to original [batch_size, seq_len, hidden_size]
244 final_tensor = final_tensor.reshape(batch_size, seq_len, hidden_size)

```

```

243         final_tensor = final_tensor.reshape(patch_size, seq_len, hidden_size,
244
245     return final_tensor, load_balancing_loss
246
247
248 class VisionLayer(nn.Module):
249     def __init__(self,
250                 hidden_dim: int,
251                 head_dim: int,
252                 q_head: int,
253                 kv_head: int,
254                 device: str,
255                 expansion_factor: int = 4,
256                 dropout_ratio: float = 0.1,
257                 use_moe: bool = False,
258                 num_experts: int = 8,
259                 lora_rank: int = 16):
260         super().__init__()
261         self.use_moe = use_moe
262         self.device = device
263
264         self.norm1 = nn.LayerNorm(hidden_dim)
265         self.attention = VisionAttention(hidden_dim, head_dim, q_head, kv_head)
266
267         self.norm2 = nn.LayerNorm(hidden_dim)
268         if self.use_moe:
269             self.moe = MOE(hidden_dim, device, num_experts=num_experts, expansion_factor=expansion_factor,
270                             dropout_ratio=dropout_ratio, lora_rank=lora_rank)
271         else:
272             self.ffn = FeedForward(hidden_dim, expansion_factor=expansion_factor, dropout_ratio=dropout_ratio,
273                                     lora_rank=lora_rank)
274
275     def forward(self, tensor: torch.Tensor, attention_mask: torch.Tensor = None, skip_connection: torch.Tensor = None,
276                 fine_tuning: bool = False):
277         skip_connection = tensor
278         tensor = self.norm1(tensor)
279         tensor = self.attention(tensor, attention_mask=attention_mask, fine_tuning=fine_tuning)
280         tensor += skip_connection
281
282         skip_connection = tensor
283         tensor = self.norm2(tensor)
284         if self.use_moe:
285             tensor, load_balancing_loss = self.moe(tensor, fine_tuning=fine_tuning)
286         else:
287             tensor = self.ffn(tensor, fine_tuning=fine_tuning)
288             load_balancing_loss = torch.tensor(0.0, dtype=tensor.dtype, device=tensor.device)

```

```
289         tensor += skip_connection
290
291     return tensor, load_balancing_loss
292
293
294 class VisionTransformer(nn.Module):
295     def __init__(self,
296                 image_size: tuple,
297                 num_classes: int = 1,
298                 patch_size: int = 8,
299                 in_channels: int = 3,
300                 num_layer: int = 3,
301                 hidden_dim: int = 1024,
302                 expansion_factor: int = 8,
303                 head_dim: int = 64,
304                 q_head: int = 16,
305                 kv_head: int = 4,
306                 dropout_ratio: float = 0.1,
307                 use_moe: bool = True,
308                 num_experts: int = 8,
309                 load_balancing_loss_weight: float = 1e-2,
310                 fine_tuning: bool = False,
311                 lora_rank: int = 16):
312         super().__init__()
313         self.device = device
314         self.num_layer = num_layer
315         self.load_balancing_loss_weight = load_balancing_loss_weight
316         self.fine_tuning = fine_tuning
317
318         # Patch embedding
319         self.patch_embedding = PatchEmbedding(
320             image_size=image_size,
321             patch_size=patch_size,
322             in_channels=in_channels,
323             embedding_dim=hidden_dim
324         )
325
326         # Calculate number of patches (sequence length)
327         self.num_patches = (image_size[0] // patch_size) * (image_size[1] // patch_size)
328
329         if q_head == None:
330             q_head = (hidden_dim // head_dim)
331
332         if kv_head == None:
333             kv_head = (hidden_dim // head_dim)
```



```

334
335     if hidden_dim % (head_dim * q_head) != 0 or hidden_dim % (head_dim * l
336         raise ValueError("Error: hidden_dim or projection_dim (if specific
337
338     # Create transformer layers
339     self.transformer = nn.ModuleList()
340     for _ in range(self.num_layer):
341         self.transformer.append(VisionLayer(
342             hidden_dim, head_dim, q_head, kv_head, device,
343             expansion_factor=expansion_factor,
344             dropout_ratio=dropout_ratio,
345             use_moe=use_moe,
346             num_experts=num_experts,
347             lora_rank=lora_rank
348         ))
349     self.output_norm = nn.LayerNorm(hidden_dim)
350
351     # Final classifier head
352     self.classifier = nn.Linear(hidden_dim, num_classes)
353
354     def begin_fine_tunning(self) -> None:
355         self.fine_tuning = True
356         for name, param in self.named_parameters():
357             if "lora" not in name:
358                 param.requires_grad = False
359             else:
360                 param.requires_grad = True
361
362     def exit_fine_tunning(self) -> None:
363         self.fine_tuning = False
364         for name, param in self.named_parameters():
365             if "positions" in name:
366                 param.requires_grad = False
367             else:
368                 param.requires_grad = True
369
370     def forward(self, x: torch.Tensor) -> tuple[torch.Tensor, torch.Tensor]:
371         # Handle input shape
372         if len(x.shape) == 3: # [batch_size, 64, 72]
373             # Reshape to [batch_size, channels, height, width]
374             # Assuming the input is grayscale (1 channel)
375             batch_size, height, width = x.shape
376             x = x.unsqueeze(1) # Add channel dimension [batch_size, 1, 64, 72]
377
378         # Apply patch embedding

```

```

379         x = self.patch_embedding(x)
380
381         # Track load-balancing across layers (only if MoE is used)
382         load_balancing_sum = torch.tensor(0.0, device=self.device)
383
384         # Pass through transformer layers
385         for layer in self.transformer:
386             x, load_balancing_loss = layer(x, fine_tuning=self.fine_tuning)
387             load_balancing_sum += load_balancing_loss
388
389         load_balancing_loss = (load_balancing_sum / self.num_layer) * self.log
390
391         # Apply output normalization
392         x = self.output_norm(x)
393
394         # Use CLS token for classification
395         x = x[:, 0] # Take only the CLS token
396
397         # Apply classifier
398         x = self.classifier(x)
399
400         return x, load_balancing_loss

```

✓ Preprocessing Data

```

1 data0 = np.load('Run357479_Dataset_iodic.npy')
2 data1 = np.load('Run355456_Dataset_jqkne.npy')
3
4 # Create labels: 0 for class 0 and 1 for class 1
5 labels_0 = np.zeros((data0.shape[0],), dtype=np.int32)
6 labels_1 = np.ones((data1.shape[0],), dtype=np.int32)
7
8 # Concatenate data and labels
9 data = np.concatenate([data0, data1], axis=0) # Shape (20000, 64, 72)
10 labels = np.concatenate([labels_0, labels_1], axis=0) # Shape (20000,)
11
12 # Shuffle data and labels together
13 indices = np.random.permutation(data.shape[0])
14 data = data[indices]
15 labels = labels[indices]
16
17 # Separate into training, validation and testing set.

```

```

18 n_total = data.shape[0]
19 n_train = int(n_total * 0.8)
20 n_val = int(n_total * 0.1)
21 n_test = n_total - n_train - n_val
22
23 train_data, val_data, test_data = data[:n_train], data[n_train:n_train + n_val], data[n_train + n_val:]
24 train_labels, val_labels, test_labels = labels[:n_train], labels[n_train:n_train + n_val], labels[n_train + n_val:]
25
26 # Compute mean and std from the training data only
27 mean = train_data.mean()
28 std = train_data.std()
29
30 # Apply the same transformation to train, val, test
31 train_data = (train_data - mean) / (std + 1e-7)
32 val_data = (val_data - mean) / (std + 1e-7)
33 test_data = (test_data - mean) / (std + 1e-7)
34
35 # Convert to PyTorch tensors
36 train_data_tensor = torch.tensor(train_data, dtype=torch.float32)
37 val_data_tensor = torch.tensor(val_data, dtype=torch.float32)
38 test_data_tensor = torch.tensor(test_data, dtype=torch.float32)
39
40 train_labels_tensor = torch.tensor(train_labels, dtype=torch.long)
41 val_labels_tensor = torch.tensor(val_labels, dtype=torch.long)
42 test_labels_tensor = torch.tensor(test_labels, dtype=torch.long)
43
44 # Create PyTorch DataLoaders
45 train_dataset = TensorDataset(train_data_tensor, train_labels_tensor.float().unsqueeze(-1))
46 val_dataset = TensorDataset(val_data_tensor, val_labels_tensor.float().unsqueeze(-1))
47 test_dataset = TensorDataset(test_data_tensor, test_labels_tensor.float().unsqueeze(-1))
48
49 train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
50 val_loader = DataLoader(val_dataset, batch_size=batch_size)
51 test_loader = DataLoader(test_dataset, batch_size=batch_size)

```

✓ Preparing for training

```
1 # Checkpoint loading
2 def load_checkpoint(model, optimizer, filepath):
3     checkpoint = torch.load(filepath)
4     model.load_state_dict(checkpoint['model_state_dict'])
5     optimizer.load_state_dict(checkpoint['optimizer_state_dict'])
6     epoch = checkpoint['epoch']
7     validation_loss = checkpoint.get('validation_loss', float('inf'))
8     print(f"Loaded checkpoint from epoch {epoch} with validation loss {validation_loss}")
9     return epoch
10
11 # Checkpoint saving
12 def save_checkpoint(model, optimizer, epoch, validation_loss):
13     checkpoint = {
14         'model_state_dict': model.state_dict(),
15         'optimizer_state_dict': optimizer.state_dict(),
16         'epoch': epoch,
17         'validation_loss': validation_loss
18     }
19     torch.save(checkpoint, f"{save_dir}/vit_checkpoint_epoch_{epoch}.pt")
20     # Save best model separately
21     if epoch == 0 or validation_loss < min(loss_valid):
22         torch.save(checkpoint, f"{save_dir}/vit_best_model.pt")
23         print(f"Saved best model with validation loss: {validation_loss:.6f}")
```

```

1 vit = VisionTransformer(
2     image_size=(64, 72),      # Your input image dimensions
3     patch_size=8,             # Size of each patch
4     in_channels=1,            # We technically don't have channel value here.
5     num_classes=1,           # Number of output classes
6     num_layer=3,              # Number of transformer layers
7     hidden_dim=256,           # Hidden dimension
8     expansion_factor=4,       # Expansion factor for FFN
9     head_dim=64,              # Dimension of each attention head
10    q_head=4,                  # Number of query heads
11    kv_head=1,                 # Number of key/value heads
12    use_moe=True,              # Whether to use Mixture of Experts
13    num_experts=4
14 ).to(device)
15
16 # Load checkpoint if available
17 current_epoch = 0
18 if checkpoint_filepath is not None and checkpoint_filepath != "":
19     current_epoch = load_checkpoint(vit, optimizer, checkpoint_filepath) + 1
20
21 print(f"This model has {sum(p.numel() for p in vit.parameters())} parameters.")
22 print(f"Training on {device}")

```

➡ This model has 10774253 parameters.
Training on cuda

```

1 # Binary classification loss since num_classes=1
2 criterion = nn.BCEWithLogitsLoss()
3 optimizer = torch.optim.AdamW(vit.parameters(), lr=lr, weight_decay=weight_decay)
4 scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=epoch)
5 scaler = amp.GradScaler() # Initialize gradient scaler for mixed precision training

```

```

1 # Initialize loss tracking lists
2 loss_train = []
3 loss_valid = []
4 accuracy_train = []
5 accuracy_valid = []

```

✓ Training

```
1 # Training loop
```

```

1 # Training loop
2 for epoch in range(current_epoch, epochs):
3     print(f"Epoch {epoch+1}/{epochs}")
4
5     # Training phase
6     vit.train()
7     loss_train_epoch = []
8     correct_train = 0
9     total_train = 0
10
11     for inputs, targets in tqdm(train_loader, desc="Training"):
12         inputs = inputs.to(device).float()
13         targets = targets.to(device).float()
14
15         # Forward pass with mixed precision
16         with amp.autocast():
17             outputs, load_balancing_loss = vit(inputs)
18             loss = criterion(outputs, targets) + load_balancing_loss
19
20         # Backward pass with gradient scaling
21         scaler.scale(loss).backward()
22
23         # Gradient clipping
24         scaler.unscale_(optimizer)
25         torch.nn.utils.clip_grad_norm_(vit.parameters(), max_norm=1.0)
26
27         # Optimizer step with scaler
28         scaler.step(optimizer)
29         scaler.update()
30
31         # Zero gradients
32         optimizer.zero_grad()
33
34         # Update scheduler
35         scheduler.step()
36
37         # Record loss
38         loss_train_epoch.append(loss.item())
39
40         # Calculate accuracy
41         predicted = (torch.sigmoid(outputs) > 0.5).float()
42         total_train += targets.size(0)
43         correct_train += (predicted == targets).sum().item()
44
45     # Calculate epoch statistics

```

```
46 epoch_loss = np.mean(loss_train_epoch)
47 epoch_accuracy = 100 * correct_train / total_train
48 loss_train.append(epoch_loss)
49 accuracy_train.append(epoch_accuracy)
50
51 # Validation phase
52 vit.eval()
53 loss_val_epoch = []
54 correct_val = 0
55 total_val = 0
56
57 with torch.no_grad():
58     for inputs, targets in tqdm(val_loader, desc="Validation"):
59         inputs = inputs.to(device).float()
60         targets = targets.to(device).float()
61
62         # Forward pass
63         with amp.autocast():
64             outputs, load_balancing_loss = vit(inputs)
65             loss = criterion(outputs, targets) + load_balancing_loss
66
67         # Record loss
68         loss_val_epoch.append(loss.item())
69
70         # Calculate accuracy
71         predicted = (torch.sigmoid(outputs) > 0.5).float()
72         total_val += targets.size(0)
73         correct_val += (predicted == targets).sum().item()
74
75 # Calculate epoch validation statistics
76 epoch_val_loss = np.mean(loss_val_epoch)
77 epoch_val_accuracy = 100 * correct_val / total_val
78 loss_valid.append(epoch_val_loss)
79 accuracy_valid.append(epoch_val_accuracy)
80
81 # Print epoch results
82 print(f"Training - Loss: {epoch_loss:.6f}, Accuracy: {epoch_accuracy:.2f}")
83 print(f"Validation - Loss: {epoch_val_loss:.6f}, Accuracy: {epoch_val_ac")
84
85 # Plot and save training progress
86 plt.figure(figsize=(12, 5))
87
88 plt.subplot(1, 2, 1)
89 plt.plot(loss_train, label="Training loss")
90 plt.plot(loss_valid, label="Validation loss")
```

```

91     plt.xlabel("Epoch")
92     plt.ylabel("Loss")
93     plt.legend()
94     plt.title("Training and Validation Loss")
95
96     plt.subplot(1, 2, 2)
97     plt.plot(accuracy_train, label="Training accuracy")
98     plt.plot(accuracy_valid, label="Validation accuracy")
99     plt.xlabel("Epoch")
100    plt.ylabel("Accuracy (%)")
101    plt.legend()
102    plt.title("Training and Validation Accuracy")
103
104    plt.tight_layout()
105    plt.savefig(f"{save_dir}/training_progress_epoch_{epoch}.png")
106    plt.close()
107
108 # Save checkpoint
109 save_checkpoint(vit, optimizer, epoch, epoch_val_loss)

```

```

➡ Epoch 1/5
Training: 100%|██████████| 16/16 [00:06<00:00, 2.63it/s]
Validation: 100%|██████████| 2/2 [00:00<00:00, 3.58it/s]
Training - Loss: 0.579171, Accuracy: 74.58%
Validation - Loss: 0.027781, Accuracy: 99.70%
Epoch 2/5
Training: 100%|██████████| 16/16 [00:04<00:00, 3.26it/s]
Validation: 100%|██████████| 2/2 [00:00<00:00, 3.17it/s]
Training - Loss: 0.016199, Accuracy: 99.88%
Validation - Loss: 0.010464, Accuracy: 100.00%
Epoch 3/5
Training: 100%|██████████| 16/16 [00:04<00:00, 3.29it/s]
Validation: 100%|██████████| 2/2 [00:00<00:00, 3.11it/s]
Training - Loss: 0.010416, Accuracy: 100.00%
Validation - Loss: 0.010539, Accuracy: 100.00%
Epoch 4/5
Training: 100%|██████████| 16/16 [00:04<00:00, 3.20it/s]
Validation: 100%|██████████| 2/2 [00:00<00:00, 3.57it/s]
Training - Loss: 0.010518, Accuracy: 100.00%
Validation - Loss: 0.010374, Accuracy: 100.00%
Epoch 5/5
Training: 100%|██████████| 16/16 [00:05<00:00, 3.15it/s]
Validation: 100%|██████████| 2/2 [00:00<00:00, 3.07it/s]
Training - Loss: 0.010164, Accuracy: 100.00%
Validation - Loss: 0.010069, Accuracy: 100.00%

```

```
1 import numpy as np
```



```
2 import matplotlib.pyplot as plt
3 from sklearn.metrics import roc_curve, auc
4
5 # Lists to store true labels and predicted probabilities
6 all_targets = []
7 all_probs = []
8
9 print("\nEvaluating on test set...")
10 vit.eval()
11 test_loss = 0
12 correct = 0
13 total = 0
14
15 with torch.no_grad():
16     for inputs, targets in tqdm(test_loader, desc="Testing"):
17         inputs = inputs.to(device).float()
18         targets = targets.to(device).float()
19
20         with amp.autocast():
21             outputs, load_balancing_loss = vit(inputs)
22             loss = criterion(outputs, targets) + load_balancing_loss
23
24             test_loss += loss.item()
25             # Compute probabilities using sigmoid
26             probabilities = torch.sigmoid(outputs)
27             # Calculate binary predictions for accuracy
28             predicted = (probabilities > 0.5).float()
29             total += targets.size(0)
30             correct += (predicted == targets).sum().item()
31
32             # Append to lists (move to CPU and convert to numpy)
33             all_targets.append(targets.cpu().numpy())
34             all_probs.append(probabilities.cpu().numpy())
35
36 avg_test_loss = test_loss / len(test_loader)
37 test_accuracy = 100 * correct / total
38
39 print(f"Test set - Loss: {avg_test_loss:.6f}, Accuracy: {test_accuracy:.2f}%")
40
41 # Concatenate all the batches
42 all_targets = np.concatenate(all_targets)
43 all_probs = np.concatenate(all_probs)
44
45 # Compute ROC curve and AUC
46 fpr, tpr, thresholds = roc_curve(all_targets, all_probs)
```


```
47 roc_auc = auc(fpr, tpr)
48 print("AUC: {:.4f}".format(roc_auc))
49
50 # Plot ROC curve
51 plt.figure()
52 plt.plot(fpr, tpr, lw=2, label='ROC curve (AUC = %0.2f)' % roc_auc)
53 plt.plot([0, 1], [0, 1], lw=2, linestyle='--', color='gray')
54 plt.xlim([0.0, 1.0])
55 plt.ylim([0.0, 1.05])
56 plt.xlabel('False Positive Rate')
57 plt.ylabel('True Positive Rate')
58 plt.title('Receiver Operating Characteristic')
59 plt.legend(loc="lower right")
60 plt.show()
61
62 # Save final model
63 torch.save({
64     'model_state_dict': vit.state_dict(),
65     'test_accuracy': test_accuracy,
66     'test_loss': avg_test_loss
67 }, f"{save_dir}/vit_final_model.pt")
68
69 print("Test completed!")
70
```



Evaluating on test set...

Testing: 100%|██████████| 2/2 [00:00<00:00, 2.61it/s]Test set - Loss: 0.01006

AUC: 1.0000



Test completed!

```
1 # Code for visualization of training data, not used in actual training but ke
2 non_zero_values = data1[data1 != 0]
3
4 # Calculate statistics
5 min_non_zero = np.min(non_zero_values)
6 max_non_zero = np.max(non_zero_values)
7 mean_non_zero = np.mean(non_zero_values)
8 median_non_zero = np.median(non_zero_values)
9 std_non_zero = np.std(non_zero_values)
10
11 # Create figure with subplots
12 fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(10, 12))
13
```

```
14 # Histogram with KDE
15 sns.histplot(non_zero_values, kde=True, ax=ax1)
16 ax1.set_title('Distribution of Non-Zero Values', fontsize=14)
17 ax1.set_xlabel('Value', fontsize=12)
18 ax1.set_ylabel('Frequency', fontsize=12)
19 ax1.axvline(min_non_zero, color='r', linestyle='--', label=f'Min: {min_non_ze
20 ax1.axvline(max_non_zero, color='g', linestyle='--', label=f'Max: {max_non_ze
21 ax1.axvline(mean_non_zero, color='b', linestyle='-', label=f'Mean: {mean_non_
22 ax1.axvline(median_non_zero, color='purple', linestyle='-.', label=f'Median:
23 ax1.legend()
24
25 # Boxplot
26 sns.boxplot(x=non_zero_values, ax=ax2)
27 ax2.set_title('Boxplot of Non-Zero Values', fontsize=14)
28 ax2.set_xlabel('Value', fontsize=12)
29
30 # Add text with statistics
31 stats_text = (f"Non-zero count: {len(non_zero_values)}\n"
32               f"Min: {min_non_zero:.2f}\n"
33               f"Max: {max_non_zero:.2f}\n"
34               f"Mean: {mean_non_zero:.2f}\n"
35               f"Median: {median_non_zero:.2f}\n"
36               f"Std Dev: {std_non_zero:.2f}")
37
38 fig.text(0.15, 0.01, stats_text, fontsize=12, bbox=dict(facecolor='white', al
39
40 plt.tight_layout(rect=[0, 0.03, 1, 0.97])
41 plt.savefig('non_zero_distribution.png')
42 plt.show()
43
44 # Print summary to console
45 print(f"Non-zero value summary:")
46 print(f"Count: {len(non_zero_values)}")
47 print(f"Min: {min_non_zero:.2f}")
48 print(f"Max: {max_non_zero:.2f}")
49 print(f"Mean: {mean_non_zero:.2f}")
50 print(f"Median: {median_non_zero:.2f}")
51 print(f"Standard deviation: {np.std(non_zero_values):.2f}")
```

